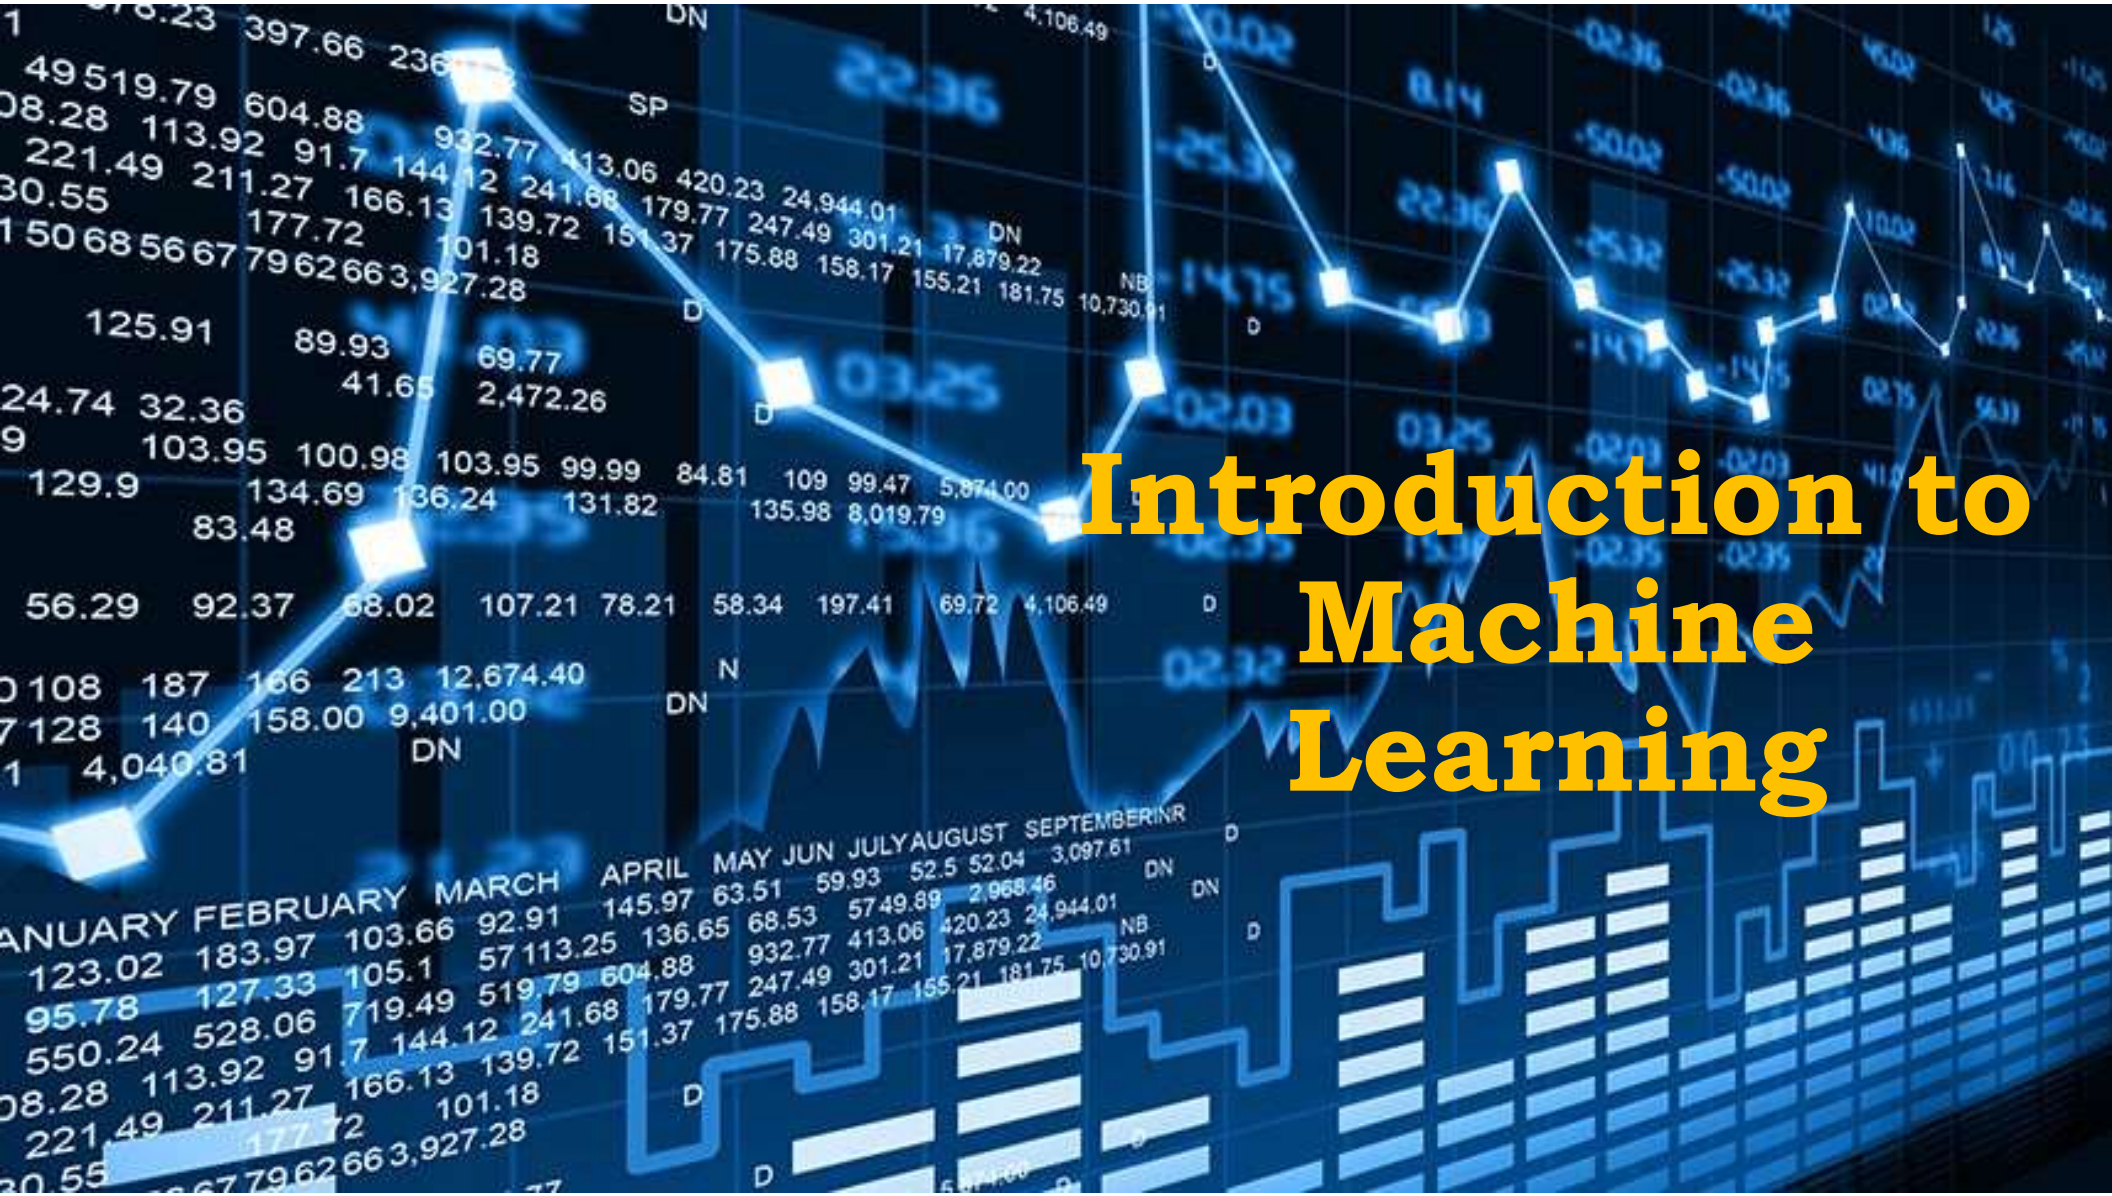
Introduction to
Machine
Learning

Data Scientist Core Skills Surrounded By Tool Skills

Machine Learning

- Definition: The use and development of computer systems that are able to learn and adapt without following explicit instructions, by using algorithms and statistical models to analyze and draw inferences from patterns in data
- In short, machine learning allows computers to learn to recognize patterns and infer/predict from data i.e. being explicitly programmed where to look!
- Let's understand this with a simple example:

Could we make computers do the guess?

Machine Learning in Daily Lives

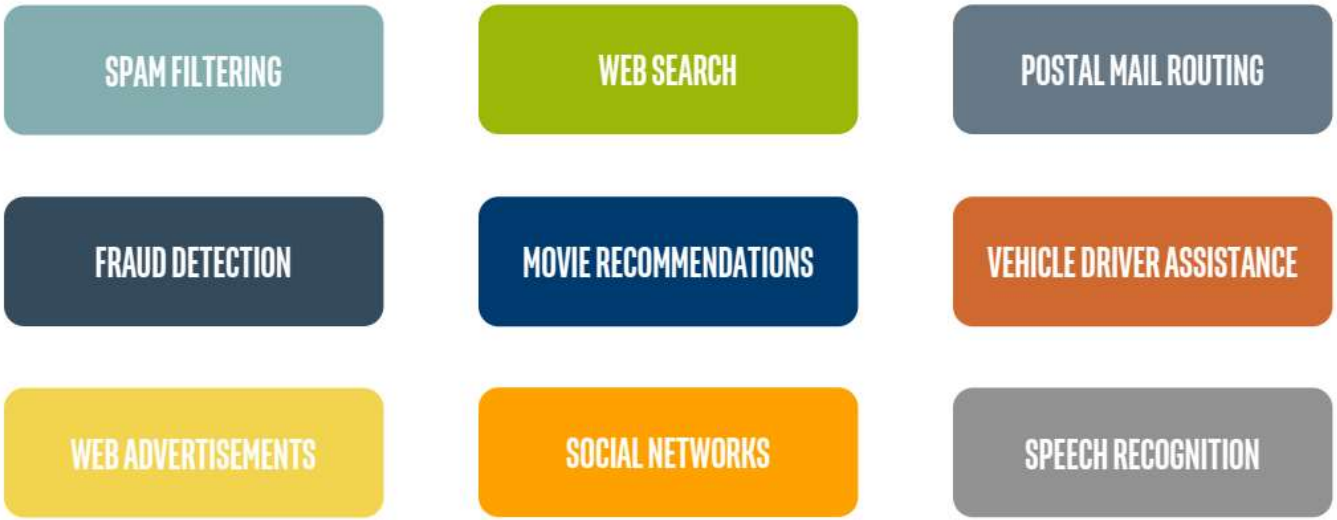

Machine Learning Vocabulary

Dataset: Tabulated data

Features: properties of the data

Example: one row

Target: column to predict

Label: target value for single data point

Sepal length	Sepal width	Petal length	Petal width	Species
6.7	3.0	5.2	2.3	Virginica
6.4	2.8	5.6	2.1	Virginica
4.6	3.4	1.4	0.3	Setosa
6.9	3.1	4.9	1.5	Versicolor
4.4	2.9	1.4	0.2	Setosa
4.8	3.0	1.4	0.1	Setosa
5.9	3.0	5.1	1.8	Virginica
5.4	3.9	1.3	0.4	Setosa
4.9	3.0	1.4	0.2	Setosa
5.4	3.4	1.7	0.2	Setosa

Machine Learning Categories

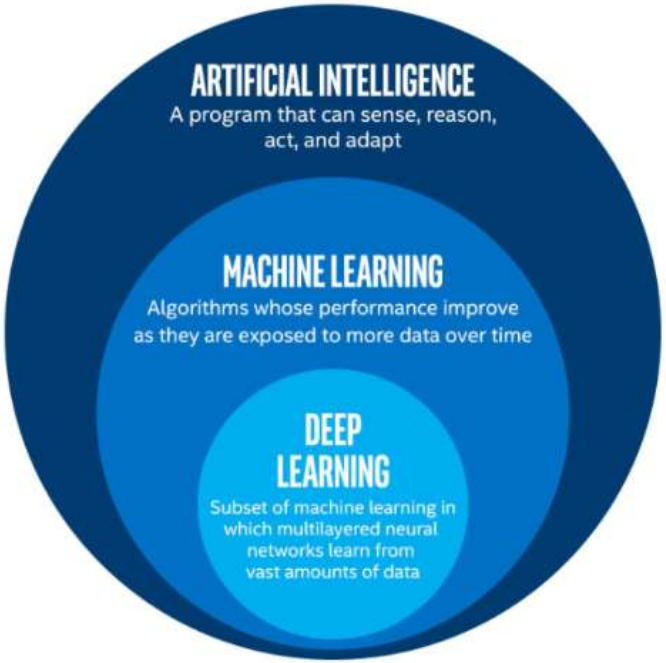

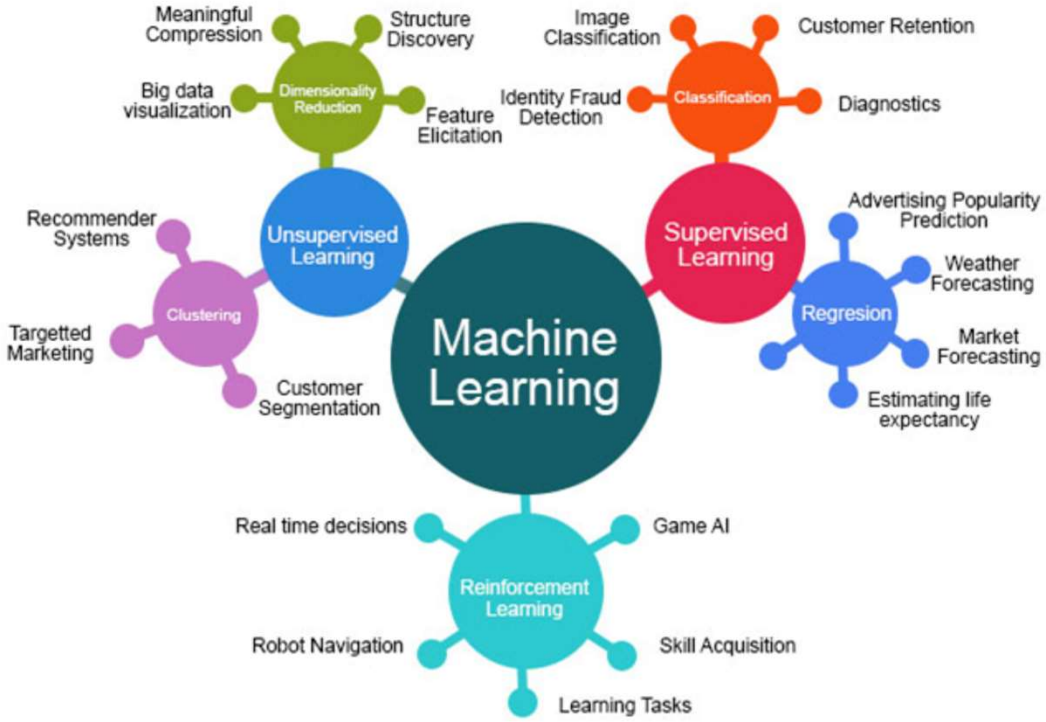

Supervised Learning

- For each observation of the predictor measurement(s) xi, i = 1, . . . , n there is an associated response measurement yi.
- We wish to fit a model that relates the response to the predictors, with two aims:
  - Accurately predicting the response for future observations (prediction)
  - Better understanding the relationship between the response and the predictors (inference).
- We could use the supervised learning algorithms to the following scenarios:
  - Anticipate which credit card transactions could be fraudulent
  - Which insurance customer is likely to file a claim
  - Predict the price of the house based on different features

Sample Data for Supervised Learning

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa
5	6	5.4	3.9	1.7	0.4	Iris-setosa
6	7	4.6	3.4	1.4	0.3	Iris-setosa
7	8	5.0	3.4	1.5	0.2	Iris-setosa
8	9	4.4	2.9	1.4	0.2	Iris-setosa
9	10	4.9	3.1	1.5	0.1	Iris-setosa

Unsupervised Learning

- Unsupervised learning describes the somewhat more challenging situation in which for every observation i = 1, . . . , n, we observe a vector of measurements xi but no associated response yi.
- It is not possible to fit a linear regression model, since there is no response variable to predict.
- In this setting, we are in some sense working blind; the situation is referred to as unsupervised because we lack a response variable that can supervise our analysis.
- The goal is to find patterns within the data
- Can be used in scenarios such as:
  - Text segmentation
  - Recommending items

Sample Data for Un-supervised Learning

Country	RedMeat	WhiteMeat	Eggs	Milk	Fish	Cereals	Starch	Nuts	Fr&Veg
Albania	10.1	1.4	0.5	8.9	0.2	42.3	0.6	5.5	1.7
Austria	8.9	14	4.3	19.9	2.1	28	3.6	1.3	4.3
Belgium	13.5	9.3	4.1	17.5	4.5	26.6	5.7	2.1	4
Bulgaria	7.8	6	1.6	8.3	1.2	56.7	1.1	3.7	4.2
Czechoslovakia	9.7	11.4	2.8	12.5	2	34.3	5	1.1	4
Denmark	10.6	10.8	3.7	25	9.9	21.9	4.8	0.7	2.4
E Germany	8.4	11.6	3.7	11.1	5.4	24.6	6.5	0.8	3.6
Finland	9.5	4.9	2.7	33.7	5.8	26.3	5.1	1	1.4
France	18	9.9	3.3	19.5	5.7	28.1	4.8	2.4	6.5

Reinforcement Learning

- Often used for robotics, gaming and navigation
- With reinforcement learning, the algorithm discovers through trial and error which actions yield the greatest rewards.
- Has three primary components:
  - Agent (the learner or decision maker)
  - Environment (everything the agent interacts with)
  - Actions (what agent can do)
- The objective is for the agent to choose actions that yield expected reward over a given amount of time
- The agent will reach the goal using a good policy, there for a goal is to figure out a good policy

Machine Learning Process